Events (ID, event_date, start_time, end_time, home_team, visiting_team, created_at)
PK: ID
FK: home_team ◊ Teams.ID, visiting_team ◊ Teams.ID

Teams (ID, team_name, team_desc, country, *sport_id*, created_at)
PK: ID
FK: sport_id ◊ Sports.ID

Sports (ID, sport_name, sport_desc, created_at)
PK: ID

The logical schema above, unlike the ER diagram, includes referential integrity constraints as well as timestamps and primary key constraints (ID) but does not include relationship mappings. Primary keys are underlined and foreign keys written cursive.

Nevertheless, worth mentioning is that although the **Events** table has 'ID' as its primary key, two other composite candidate keys are primarily used to prevent the insertion of duplicate rows, namely **(home_team, event_date)** and **(visiting_team, event_date)**.

The tables could be further decomposed; first normal form was violated in the 'teams' relation whenever a team was from the US – the country attribute was non-atomic (e.g. country: California, USA) , but due to the scale of the project this was not relevant.